

Balaji Chidambaram

(669) 263-5630 | San Francisco Bay Area, CA | balaji21.chidambaram@gmail.com | linkedin.com/in/balaji-chidam | balachidam21.github.io

SUMMARY

Software engineer building and operating production agentic and GenAI systems - multi-agent orchestration and RAG at enterprise scale on AWS, Databricks and Google Cloud, with the policy-scoped tool access, audit lineage and automated evaluation that keep them trustworthy. MS in Computer Science, USC.

EXPERIENCE

Genpact — Principal Consultant, AI Engineer

Mar 2026 - Present

Remote, USA

- Lead architect for an enterprise Agentic xP&A platform unifying Financial Planning, Supply Chain, Workforce Planning and Demand Planning across 4 business segments. Ran client discovery, set the multi-agent reference architecture, and shipped a production MVP on Databricks in 1 month with a cross-functional team of 3 engineers and 2 SMEs.
- Architected a Lakebase-backed agent governance layer where finance users curate 400+ KPIs and drivers directly, with versioned lineage - clone, edit, publish, audit. Cut data-readiness prep from over a month to a few days, with every agent output traceable to a governed source.

Genpact — Assistant Manager, AI Engineer

May 2024 - Feb 2026

Remote, USA

- Architected a multi-agent Conversational AI platform using a LangGraph orchestrator over 8 deterministic tools, reducing analysis lead time from days to minutes for 100+ enterprise users.
- Exposed 20+ reusable analysis skills through MCP servers with least-privilege tool access and prompt-injection guardrails on every tool call.
- Systematized agent quality assurance with an MLflow GenAI evaluation harness using LLM-as-a-judge scorers - raising user acceptance from 40% to 60% and preempting ~20 pre-production regressions.
- Engineered a query-lineage governance layer surfacing underlying SQL and source datasets, moving AI-driven analysis from a black box to an auditable solution.

Keck Medicine of USC — Data Engineer (Contract)

Feb 2024 - Dec 2024

Los Angeles, CA

- Engineered a multi-source data warehouse and analytics platform on AWS Redshift, integrating Salesforce Data Cloud, Google Analytics and CRM to unify analytics across data silos.
- Engineered customer segments and activations in Salesforce Data Cloud using tailored AWS pipelines (Lambda, RDS), driving a 30% increase in operational efficiency.

Keck Medicine of USC — Marketing Data Analyst

Sep 2022 - Dec 2023

Los Angeles, CA

- Built AWS-based ingestion and ETL transformation pipelines (Lambda, EC2, VPC, RDS) that cut reporting prep time by 40%. Revamped data architecture with Adverity to align financial and marketing data, improving ROI-validation accuracy by 20%.
- Designed a Tableau/AWS/Looker analytics platform on the data lake, raising campaign visibility by 54% and engagement by 32%.

Sayari Labs — Data Engineer Intern

May 2023 - Jul 2023

Washington, D.C.

- Built scalable microservice-style data pipelines (Scrapy, Kubernetes, Airflow) ingesting and transforming financial and corporate data from 10+ sources monthly with high availability.
- Standardized Spark data modeling and developed end-to-end pipeline frameworks with component-level unit tests, CI/CD integration and documentation, enabling agile cross-functional delivery.

EDUCATION

M.S. Computer Science — University of Southern California

Jan 2022 - Dec 2023

Los Angeles, CA | GPA 3.81 / 4.0

B.E. Computer Science and Engineering — Anna University

Aug 2017 - Jul 2021

India | GPA 8.97 / 10

PUBLICATION

P. Shekhar, R. Saha, M. J. B. Dudekula, B. Chidambaram. "Effectiveness of Retrieval Augmented Generation, Contextualized Examples and Prompt Finetuning on Data Enrichment, Cleaning and Master Data Creation." *IEEE ACDSA*, 2026. DOI: [10.1109/ACDSA67686.2026.11467982](https://doi.org/10.1109/ACDSA67686.2026.11467982)

SKILLS

- **Agentic AI & LLM systems:** Multi-agent systems, ReAct, Hierarchical delegation, Supervisor pattern, MCP, Agent Skills, RAG, Vector databases (pgvector)
- **Languages:** Python, PySpark, SQL, C++, C, Java, React
- **AI / ML frameworks:** LangGraph, MCP servers, Google ADK, TensorFlow, PyTorch, MLflow, TimeGPT, Transformers
- **Data & streaming:** Apache Spark, Kafka, Cassandra, Airflow, Redshift, Salesforce Data Cloud, Scrapy
- **Cloud & platform:** AWS (Lambda, EC2, VPC, RDS, Redshift), Google Cloud (Vertex AI), Databricks, Databricks Apps, Lakebase, Azure
- **Analytics & interfaces:** Power BI, Tableau, Looker, Streamlit
- **Certifications:** [Google Cloud Professional Machine Learning Engineer](#), [Databricks Generative AI Associate](#), [Google Cloud Associate Cloud Engineer](#)